

*8*

# Overview

1. Exploratory Data Analysis (EDA): what and why?
2. Things to explore
3. Exploration and visualization tools
4. (A bit of) dataset cleaning
5. Kaggle competition EDA

## Check if the data is intuitive

...	Age	...
...	21	...
...	45	...
...	336	...
...	19	...
...	...	...

- Is 336 a typo?
- Or we misinterpret the feature and age 336 is normal

# Motivating example

id	...	# promos sent	# promos used	diff	used this promo?
13	...	0	0	1	1
13	...	1	1	0	0
13	...	2	1	1	0
13	...	4	2	1	1
13	...	5	3	1	1
13	...	6	3	NaN	0

1. For each person sort by '# promos sent'
2. Look at difference between consecutive rows in '# promos used' column ('diff' feature)

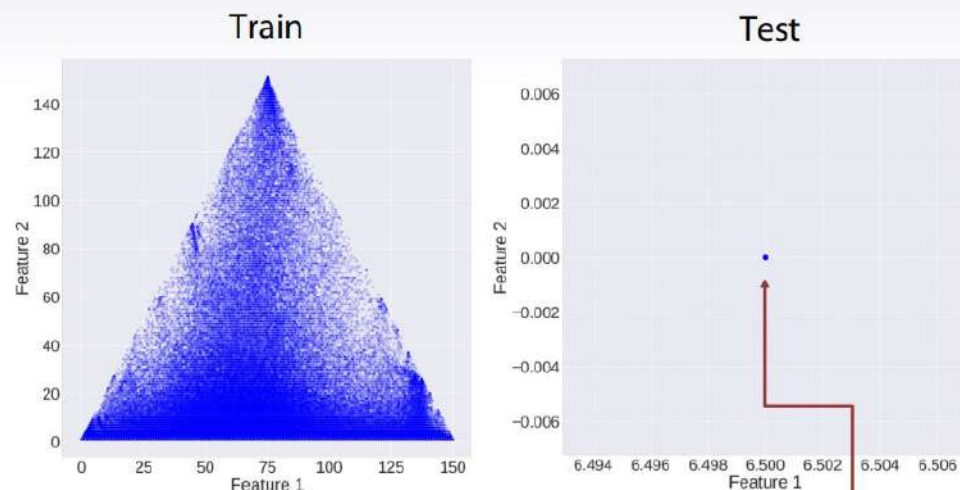
## Check if the data is intuitive

Task: Predict advertiser's cost

Data:

AdGroupId	AdNetwork Type2	MaxCpc	Slot	Clicks	Impressions	is_incorrect
78db044136	s	0.28	s_2	3	0	True
68a0110c33	s	1	s_2	1	13	False
2r39fw11w3	p	1.2	p_1	3	419	False

# Understand how the data was generated



#days in *train* > #days in *test*  
#rows in *train* < #rows in *test*

# Conclusion

- **Get domain knowledge**
  - It helps to deeper understand the problem
- **Check if the data is intuitive**
  - And agrees with domain knowledge
- **Understand how the data was generated**
  - As it is crucial to set up a proper validation

## Anonymized data

Text	Encoded text
I want this table	7ugy 972h 98ww hj34
Table is what I want	hj34 4f08 rtte 7ugy 972h
This table is red	98ww hj34 4f08 4rj9
And this is me	jk8r 98ww 4f08 9jo4

- Two things to do with anonymized features:
  - **Try to decode the features**
    - Guess the true meaning of the feature
  - **Guess the feature types**
    - Each type needs its own preprocessing


## Anonymized data

id	x1	x2	x3	x4	x5	x6
1	m268i97y	0	NO	105.4	14	
2	j0gheu6	1	YES	25.631	12	
3	26fmsp6u	1	NO	12.0	12	m268i97y
4	13e5dpzp	0	NO	140.12	14	m268i97y

- **Explore individual features**
  - Guess the meaning of the columns
  - Guess the types of the column
- **Explore feature relations**
  - Find relations between pairs
  - Find feature groups




### Histograms



```
plt.hist(x)
```


### Plot (index versus value)



```
plt.plot(x, '.')
```

Horizontal lines show repeat value, no vertical lines show shuffled ro


### Other tools



```
x.value_counts()
x.isnull()
```

It feels like 0 is there many times, but log leke pata chala asa nahi hai


### Histograms



```
plt.hist(x)
```


### Not shuffled rows

### Plot (index versus value)




```
plt.scatter(range(len(x)), x, c=y)
```

## Exploring feature relations




```
plt.scatter(x1, x2)
```




Clearly they follow  $x_2 < -x_1$ . So we should make new feature as difference between  $x_1, x_2$

## Exploring individual features: pairs



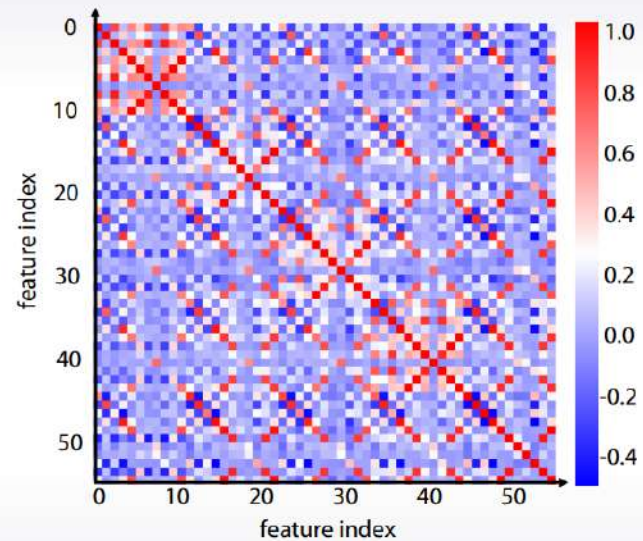
```
pd.scatter_matrix(df)
```




```
plt.scatter(x1, x2)
```

Test in right side is far away from everyone, which is bad  
We should check if we have been doing things correctly

## Exploring individual features: pairs/groups



This is same correlation matrix just after running Kmeans so we have feature groups




Tools:

```
df.corr(), plt.matshow( ... )
```



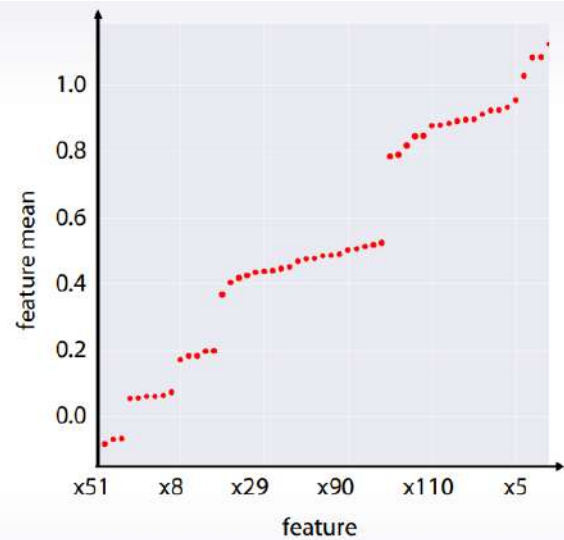
## Exploring individual features: groups



Tools:

```
df.mean().plot(style='.') 
```

Just sort karne se nice relations mil gaye



Tools:

```
df.mean().sort_values().plot(style='.') 
```

## Duplicated and constant features

is_train	f0	f1	f2	f3	f4	f5
True	13	H	1.2	1.2	A	C
True	13	H	36.6	36.6	B	A
False	13	H	0	0	A	C
False	13	G	-14	-14	C	B

```
traintest.nunique(axis=1) == 1
```

is_train	f0	f1	f2	f3	f4	f5
True	13	H	1.2	1.2	A	C
True	13	H	36.6	36.6	B	A
False	13	H	0	0	A	C
False	13	G	-14	-14	C	B

```
traintest.T.drop_duplicates()
```

is_train	f0	f1	f2	f3	f4	f5
True	13	H	1.2	1.2	A	C
True	13	H	36.6	36.6	B	A
False	13	H	0	0	A	C
False	13	G	-14	-14	C	B

```
for f in categorical_feats:
    traintest[f] = traintest[f].factorize()
traintest.T.drop_duplicates()
```

f4,f5 are same when label encoded

## EDA check list

- Get domain knowledge
  - Check if the data is intuitive
  - Understand how the data was generated
- 
- Explore individual features
  - Explore pairs and groups
- 
- Clean features up
- 
- Check for leaks! (later in this course)

Generally, overfitting refers to

- a. capturing noise
- b. capturing patterns which do not generalize to test data

In competitions, overfitting refers to

- a. low model's quality on test data, which was unexpected due to validation scores

Validation types

- Holdout: `ngroups = 1`  
`sklearn.model_selection.ShuffleSplit`
- K-fold: `ngroups = k`  
`sklearn.model_selection.Kfold`
- Leave-one-out: `ngroups = len(train)`  
`sklearn.model_selection.LeaveOneOut`

Holdout when lot of data, others when less data.  
LOO when minimal data

Stratification preserve the same target distribution over different folds

Samples and their target values

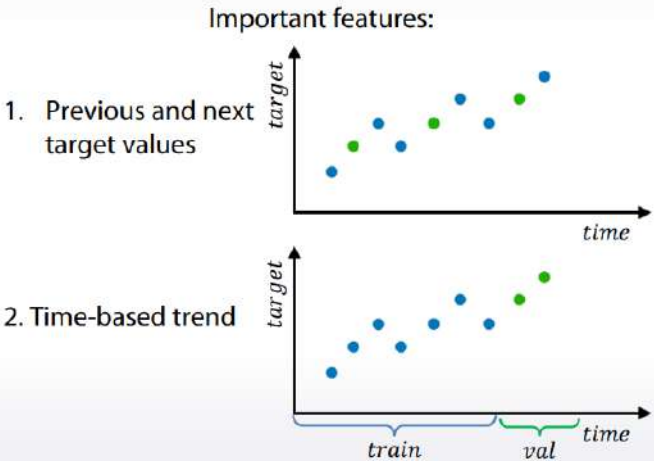
0	1	0	0	1	1	1	0
0	1	0	0	1	1	1	0
0.5		0		1		0.5	
0	1	0	0	1	1	1	0
0.5	0.5	0.5	0.5	0.5	0.5		

Stratification is useful for:

- Small datasets
- Unbalanced datasets
- Multiclass classification

Validation schemes are supposed to be used to estimate quality of the model. When you found the right hyper-parameters and want to get test predictions **don't forget to retrain your model using all training data.**

Different approaches to validation



Moving window validation

week1	week2	week3	week4	week5	week6
train			validation		
train				validation	
train					validation

1. In most cases data is split by
  - a. Row number
  - b. Time
  - c. Id
2. Logic of feature generation depends on the data splitting strategy
3. Set up your validation to mimic the train/test split of the competition




## Causes of different scores and optimal parameters

1. Too little data
2. Too diverse and inconsistent data

## We should do extensive validation


1. Average scores from different KFold splits
2. Tune model on one split, evaluate score on the other

Distribution of Heights



- Mean for train:  
Calculate from the train data
- Mean for test:  
Leaderboard probing

## Submission stage: different distributions



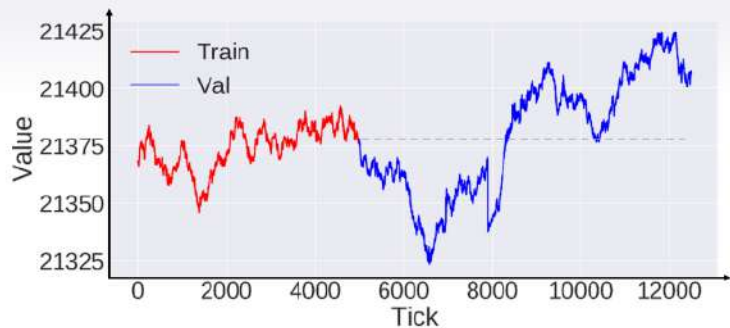
- If submission's score do not match local validation score, we should
  - Check if we have too little data in public LB
  - Check if we overfitted
  - Check if we chose correct splitting strategy
  - Check if train/test have different distributions
- Expect LB shuffle because of
  - Randomness
  - Little amount of data
  - Different public/private distributions

## Leaks in time series

- Split should be done on time.
  - In real life we don't have information from future
  - In competitions first thing to look: train/public/private split, is it on time?
- Even when split by time, features may contain information about future.
  - User history in CTR tasks
  - Weather

## Unexpected information

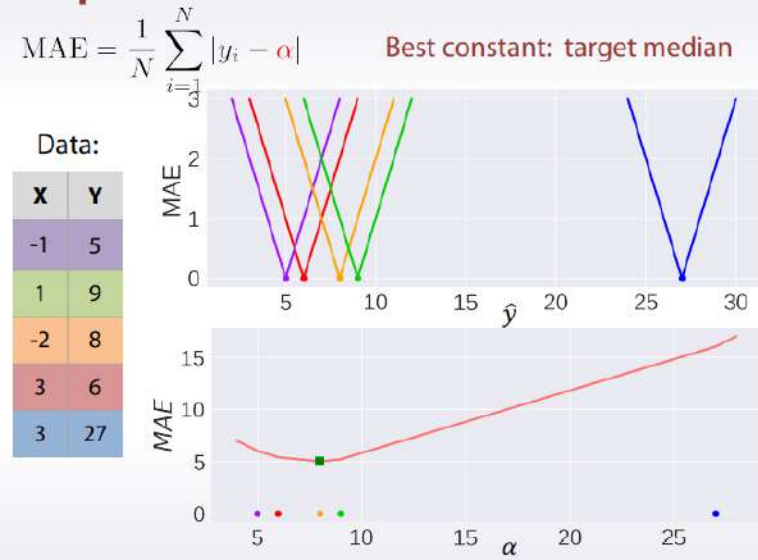
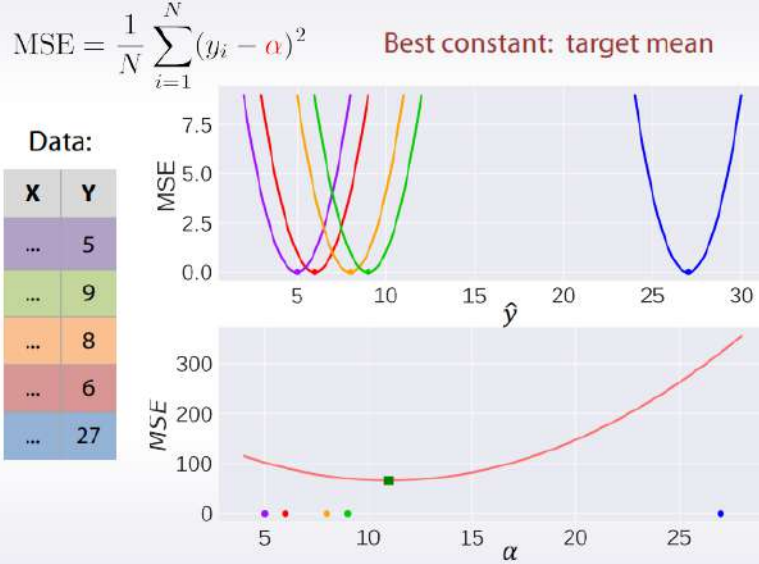
- Meta data
- Information in IDs
- Row order



$$Loss(\hat{y}_i; y_i) = \begin{cases} |y_i - \hat{y}_i|, & \text{if trend predicted correctly} \\ (y_i - \hat{y}_i)^2, & \text{if trend predicted incorrectly} \end{cases}$$

**Predict trend instead of the values:**

Predict  $y_{last} + 10^{-6}$   
or  $y_{last} - 10^{-6}$



$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

**R-squared:**

$$RMSE = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2} = \sqrt{MSE}$$

$$MSE(a) > MSE(b) \iff RMSE(a) > RMSE(b)$$

$$\frac{\partial RMSE}{\partial \hat{y}_i} = \frac{1}{2\sqrt{MSE}} \frac{\partial MSE}{\partial \hat{y}_i}$$

$$R^2 = 1 - \frac{\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2}{\frac{1}{N} \sum_{i=1}^N (y_i - \bar{y})^2} = 1 - \frac{MSE}{\frac{1}{N} \sum_{i=1}^N (y_i - \bar{y})^2}$$


$$\bar{y} = \frac{1}{N} \sum_{i=1}^N y_i$$

- **Do you have outliers in the data?**
  - Use MAE
- **Are you sure they are outliers?**
  - Use MAE
- **Or they are just unexpected values we should still care about?**
  - Use MSE
- **MSE, RMSE, R-squared**
  - They are the same from optimization perspective
- **MAE**
  - Robust to outliers




$$\text{MSPE} = \frac{100\%}{N} \sum_{i=1}^N \left( \frac{y_i - \alpha}{y_i} \right)^2$$

Best constant:  
weighted target mean



$$\text{MAPE} = \frac{100\%}{N} \sum_{i=1}^N \left| \frac{y_i - \alpha}{y_i} \right|$$

Best constant:  
weighted target median



## (R)MSLE: Root Mean Square Logarithmic Error

$$\begin{aligned} \text{RMSLE} &= \sqrt{\frac{1}{N} \sum_{i=1}^N (\log(y_i + 1) - \log(\hat{y}_i + 1))^2} = \\ &= \text{RMSE}(\log(y_i + 1), \log(\hat{y}_i + 1)) = \\ &= \sqrt{\text{MSE}(\log(y_i + 1), \log(\hat{y}_i + 1))} \end{aligned}$$

- **(R)MSPE**
  - Weighted version of MSE
- **MAPE**
  - Weighted version of MAE
- **(R)MSLE**
  - MSE in log space

$$\text{Accuracy} = \frac{1}{N} \sum_{i=1}^N [\alpha = y_i]$$

- How frequently our class prediction is correct.
- Best constant:
  - **predict the most frequent class.**

- Dataset:
  - 10 cats
  - 90 dogs

Predict always dog:  
Accuracy = **0.9!**

- Binary:

$$\text{LogLoss} = -\frac{1}{N} \sum_{i=1}^N y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)$$

$y_i \in \mathbb{R}, \hat{y}_i \in \mathbb{R}$


- Multiclass:

$$\text{LogLoss} = -\frac{1}{N} \sum_{i=1}^N \sum_{l=1}^L y_{il} \log(\hat{y}_{il})$$


$y_i \in \mathbb{R}^L, \hat{y}_i \in \mathbb{R}^L$

- In practice:

$$\text{LogLoss} = -\frac{1}{N} \sum_{i=1}^N \sum_{l=1}^L y_{il} \log(\min(\max(\hat{y}_{il}, 10^{-15}), 1 - 10^{-15}))$$



- Logloss strongly penalizes completely wrong answers



**TP** – true positives, **FP** – false positives

$$\begin{aligned} \text{AUC} &= \frac{\# \text{ correctly ordered pairs}}{\text{total number of pairs}} = \\ &= 1 - \frac{\# \text{ incorrectly ordered pairs}}{\text{total number of pairs}} \end{aligned}$$



pair = (red object, green object)

## Cohen's Kappa motivation

**Dataset:**

- 10 cats
- 90 dogs

Predict 20 cats and 80 dogs at  
random: accuracy ~ 0.74

$$0.2 \cdot 0.1 + 0.8 \cdot 0.9 = 0.74$$

$$\text{Cohen's Kappa} = 1 - \frac{1 - \text{accuracy}}{1 - p_e}$$

$p_e$  – what accuracy would be on average, if we randomly permute our predictions

$$p_e = \frac{1}{N^2} \sum_k n_{k1} n_{k2}$$

Confusion matrix  $C$ 

pred\ true	cat	dog	tiger
cat	4	2	3
dog	2	88	4
tiger	4	10	12

Weight matrix  $W$ 

pred\ true	cat	dog	tiger
cat	0	1	10
dog	1	0	10
tiger	1	1	0

Linear weights

pred\ true	1	2	3
1	0	1	2
2	1	0	1
3	2	1	0

Quadratic weights

pred\ true	1	2	3
1	0	1	4
2	1	0	1
3	4	1	0

$$\text{weighted error} = \frac{1}{\text{const}} \sum_{i,j} C_{ij} W_{ij}$$

$$\text{weighted kappa} = 1 - \frac{\text{weighted error}}{\text{weighted baseline error}}$$


## Custom loss for XGBoost

- **Define an 'objective':**
  - function that computes derivatives w.r.t. predict

```
def logregobj(pred, labels = dtrain.get_label(), preds = 1.0 / 2.0):
    grad = preds - labels
    hess = preds * (1 - preds)
    return grad, hess
```

## Early stopping

- Optimize metric **M1**, monitor metric **M2**
  - Stop when **M2 score** is the best



## Approaches for target metric optimization

- **Just run the right model!**
  - MSE, Logloss
- **Preprocess train and optimize another metric**
  - MSPE, MAPE, RMSLE, ...
- **Optimize another metric, postprocess predictions**
  - Accuracy, Kappa
- **Write custom loss function**
  - Any if you can




- Categorical feature  
- some city
- Binary classification

	feature	feature_label	feature_mean	target
0	Moscow	1	0.4	0
1	Moscow	1	0.4	1
2	Moscow	1	0.4	1
3	Moscow	1	0.4	0
4	Moscow	1	0.4	0
5	Tver	2	0.8	1
6	Tver	2	0.8	1
7	Tver	2	0.8	1
8	Tver	2	0.8	0
9	Klin	0	0.0	0
10	Klin	0	0.0	0
11	Tver	2	0.8	1

## Springleaf example


```
dtrain = xgb.DMatrix(train_new, label=y_tr)
dvalid = xgb.DMatrix(val_new, label=y_val)

evallist = [(dtrain, 'train'), (dvalid, 'eval')]
evals_result3 = {}
model = xgb.train(xgb_par, dtrain, 3000, evals=evallist,
                  verbose_eval=30, evals_result=evals_result3, early_stopping_rounds=50)
```



Overfit

1. Label encoding gives random order. No correlation with target
2. Mean encoding helps to separate zeros from ones




In [4]:

```
means = X_tr.groupby(col).target.mean()
train_new[col+'_mean_target'] = train_new[col].map(means)
val_new[col+'_mean_target'] = val_new[col].map(means)
```

means

```
Out[4]: VAR_1277
0.0    0.358965
1.0    0.219249
2.0    0.193671
3.0    0.191143
4.0    0.191080
5.0    0.185694
```

Means calculated for a given category and mapped to training & validation



Train


Validation

	feature	feature_label	feature_mean	target		feature	feature_label	feature_mean	target
8	Tver	2	0	0	10	Klin	0	1	1
9	Klin	0	0	0	11	Tver	2	1	1

## Regularization. CV loop

- Robust and intuitive
- Usually decent results with 4-5 folds across different datasets
- Need to be careful with extreme situations like LOO

KFold scheme



```
y_tr = df_tr['target'].values #target variable
skf = StratifiedKFold(y_tr,5, shuffle=True,random_state=123)

for tr_ind, val_ind in skf:
    X_tr, X_val = df_tr.iloc[tr_ind], df_tr.iloc[val_ind]
    for col in cols: #iterate though the columns we want to encode
        means = X_val[col].map(X_tr.groupby(col).target.mean())
        X_val[col+'_mean_target'] = means
    train_new.iloc[val_ind] = X_val

prior = df_tr['target'].mean() #global mean
train_new.fillna(prior,inplace=True) #fill NANS with global mean
```

- Perfect feature for LOO scheme
- Target variable leakage is still present even for KFold scheme

Leave-one-out

	feature	feature_mean	target
0	Moscow	0.50	0
1	Moscow	0.25	1
2	Moscow	0.25	1
3	Moscow	0.50	0
4	Moscow	0.50	0

## Regularization.Smoothing

- Alpha controls the amount of regularization
- Only works together with some other regularization method

$$\frac{mean(target) * nrows + globalmean * alpha}{nrows + alpha}$$

## Regularization. Expanding mean

- Least amount of leakage
- No hyper parameters
- Irregular encoding quality
- Built - in in CatBoost

```
cumsum = df_tr.groupby(col)['target'].cumsum() - df_tr['target']
cumcnt = df_tr.groupby(col).cumcount()
train_new[col+'_mean_target'] = cumsum/cumcnt
```

- There are a lot ways to regularize mean encodings
- Unending battle with target variable leakage
- CV loop or Expanding mean for practical tasks